

```
[2]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```



```
df3=pd.read_csv("telecom_customer_profile.csv")  
df3
```



	customer_id	full_name	age	gender	region	tenure_months	contract_type	payment_method	monthly_charges	total_charges
0	1	Rahul Gupta	56.000000	Male	NaN	17	Quarterly	NaN	2375.90	111040.30
1	2	Sneha Gupta	69.000000	Male	West	56	Monthly	UPI	2155.58	67731.65
2	3	Neha Mehta	46.000000	F	South	9	Monthly	NaN	2279.89	61524.41
3	4	Kunal Sharma	43.494716	Male	North	38	Quarterly	UPI	1941.69	116536.33
4	5	Sanjay Das	60.000000	Male	North	39	month-to-month	UPI	1295.17	49265.47
...
347995	347996	Meera Sharma	66.000000	Male	North	51	Yearly	Cash	1394.81	97634.14
347996	347997	Sanjay Das	26.000000	NaN	North	8	month-to-month	Credit Card	2005.03	93670.63
347997	347998	Riya Singh	53.000000	Male	West	25	month-to-month	UPI	796.77	65355.84
347998	347999	Aditya Iyer	30.000000	Male	North	32	Yearly	UPI	-1.00	92114.22
347999	348000	Neha Sharma	47.000000	M	South	40	Quarterly	NaN	2371.74	25593.71

348000 rows × 10 columns


```
[37]: df3=pd.read_csv("telecom_customer_profile.csv")
df3
```

	customer_id	avg_call_minutes	avg_data_gb	customer_support_calls	late_payment_count	complaints	churn
0	1	535.6	9.29	5	5	1	1
1	2	642.3	48.12	3	7	0	1
2	3	707.5	16.16	14	7	4	0
3	4	166.1	87.89	14	5	3	0
4	5	182.0	23.47	13	5	4	0
...
347995	347996	728.4	50.35	5	2	1	0
347996	347997	663.0	NaN	1	0	1	1
347997	347998	530.0	14.49	2	8	3	1
347998	347999	231.9	36.36	10	8	2	0
347999	348000	-1.0	19.80	9	8	1	0

348000 rows × 7 columns


```
sns.histplot(df4['churn_flag'],bins=2)
plt.title("Customer Churn Distribution")
```

```
Text(0.5, 1.0, 'Customer Churn Distribution')
```




```
[51]: sns.countplot(x="engagement_level",data=df_merge,color='red')
plt.title("Engagement Level Distribution")
```

```
[51]: Text(0.5, 1.0, 'Engagement Level Distribution')
```




```
sns.countplot(x="usage_level",data=df_merge,color='yellow')
plt.title("Data Usage Level Distribution")
```

```
Text(0.5, 1.0, 'Data Usage Level Distribution')
```




```
[63]: sns.countplot(x="billing_stress",hue='churn_flag',data=df_merge,color="gold")
plt.title("Billing Stress Impact")
```

```
[63]: Text(0.5, 1.0, 'Billing Stress Impact')
```



```
[64]: sns.countplot(x="support_risk_level",hue="churn_flag",color="pink",data=df_merge)
plt.title("Support Risk Impact")
```

```
[64]: Text(0.5, 1.0, 'Support Risk Impact')
```



```
[69]: monthlymean=df_merge['monthly_charges'].mean()
monthlymean
```

```
[69]: np.float64(1280.9540939367819)
```

```
[70]: monthlymedian=df_merge['monthly_charges'].median()
monthlymedian
```

```
[70]: 1288.0349999999999
```


```
[71]: stdmonthly=df_merge['monthly_charges'].std()
stdmonthly
```

```
[71]: 710.5206814645707
```

```
[72]: sns.kdeplot(df_merge['monthly_charges'],fill=True)
```

```
[72]: sns.kdeplot(df_merge['monthly_charges'],fill=True)
```

```
[72]: <Axes: xlabel='monthly_charges', ylabel='Density'>
```



```
[73]: skrew_val=df_merge['monthly_charges'].skew()
skrew_val
```

```
[73]: np.float64(-0.0549122539202492)
```

```
[74]: kurt_val=df_merge['monthly_charges'].kurt()  
kurt_val
```

```
[74]: np.float64(-1.122485097506128)
```

```
[75]: Q1=df_merge['monthly_charges'].quantile(0.25)  
Q3=df_merge['monthly_charges'].quantile(0.75)  
IQR=Q3-Q1  
  
print(f"25th Percentile (Q1): ${Q1}")  
print(f"75th Percentile (Q3): ${Q3}")  
print(f"Interquartile Range (IQR): ${IQR}")
```

```
25th Percentile (Q1): $681.905
```

```
75th Percentile (Q3): $1891.75
```

```
Interquartile Range (IQR): $1209.845
```

```
[76]: UF=Q3+(1.5*IQR)  
LF=Q1-(1.5*IQR)  
UF
```

```
[76]: np.float64(3706.5175)
```

```
[77]: Outliers=df_merge[df_merge['monthly_charges']>UF]  
  
print(Outliers[["full_name","monthly_charges"]])
```

```
Empty DataFrame
```

```
Columns: [full_name, monthly_charges]
```

```
Index: []
```

```
[78]: cc=df_merge["churn"].value_counts(normalize=True)
prob_churn = cc[1]
prob_no_churn = cc[0]

print(f"Probability of Churn: {prob_churn:.2f}")
print(f"Probability of Staying: {prob_no_churn:.2f}")
```

Probability of Churn: 0.28
Probability of Staying: 0.72

```
[79]: # Total Probability of churn
p_churn=(df_merge['churn'] == 1).mean()
p_churn
```

```
[79]: np.float64(0.28044827586206894)
```

```
[80]: yearly_churn=(df_merge['contract_type'] == "Yearly").mean()
yearly_churn
```

```
[80]: np.float64(0.25047988505747126)
```

```
[81]: probably_yearly_churn=(df_merge[df_merge['contract_type'] == 'Yearly']['churn'] == 1).mean()
print(f"prob(one year | churned) :{probably_yearly_churn:.2%}")
```

prob(one year | churned) : 28.23%

```
[82]: probably_mon_mon_churn=(df_merge[df_merge['contract_type'] == 'month-to-month']['churn'] == 1).mean()
print(f"prob(one year | churned) :{probably_mon_mon_churn:.2%}")
```

prob(one year | churned) : 27.75%

```
[68]: from sklearn.preprocessing import StandardScaler
```

```
[69]: cols=['age','tenure_months','monthly_charges','total_charges','avg_call_minutes','avg_data_gb']
scaler=StandardScaler()
df_merge[cols]=scaler.fit_transform(df_merge[cols])
```

```
[70]: df_merge
```

	age	tenure_months	monthly_charges	total_charges	avg_call_minutes	avg_data_gb	customer_support_calls	late_payment_count	complaints	churn	...	s
0	0.857134	-0.890825	1.541049	0.830611	0.605391	-1.296956		5	5	1	1	...
1	1.746417	0.988269	1.230967	-0.173455	1.058804	0.002571		3	7	0	1	...
2	0.173070	-1.276280	1.405923	-0.317363	1.335866	-1.067037		14	7	4	0	...
3	-0.032150	0.120995	0.929933	0.958030	-0.964770	1.333558		14	5	3	0	...
4	1.130760	0.169177	0.020008	-0.601574	-0.897205	-0.822393		13	5	4	0	...
...
347995	1.541198	0.747360	0.160243	0.519803	1.424679	0.077203		5	2	1	0	...
347996	-1.195059	-1.324461	1.019079	0.427913	1.146767	-1.607866		1	0	1	1	...
347997	0.651915	-0.505369	-0.681451	-0.228536	0.581594	-1.122927		2	8	3	1	...
347998	-0.921433	-0.168096	-1.802841	0.391829	-0.685159	-0.391002		10	8	2	0	...
347999	0.241476	0.217359	1.535195	-1.150379	-1.670599	-0.945217		9	8	1	0	...

```
[73]: X=df_merge.drop(columns=['churn'])  
y=df_merge['churn']
```

```
[76]: from sklearn.model_selection import train_test_split
```

```
[80]: X_train,X_test,y_train,y_test= train_test_split(X,y, test_size=0.2,random_state=42,stratify=y)
```

```
[81]: from sklearn.linear_model import LogisticRegression  
model = LogisticRegression()  
model.fit(X_train,y_train)
```

C:\Users\Hp\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:465: ConvergenceWarning:

lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression


```
[81]: ▾ LogisticRegression ⓘ ?  
LogisticRegression()
```

```
[82]: y_pred=model.predict(X_test)
```

```
[83]: y_pred
```



```
[100]: from sklearn.metrics import roc_curve,roc_auc_score  
y_prob=model.predict_proba(X_test)[:,1]  
fpr,tpr,_=roc_curve(y_test,y_prob)  
auc=roc_auc_score(y_test,y_prob)  
plt.plot(fpr,tpr,label=f"AUC={auc:.2f}")  
plt.xlabel("false positive rate")  
plt.ylabel("true positive rate")  
plt.title("ROC Curve")  
plt.legend()  
plt.show()
```



[110]:

```
coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': model.coef_[0]
}).sort_values(by='Coefficient', ascending=False)

coef_df.head(10)
```

[110]:

	Feature	Coefficient
18	usage_level_low usage	0.030324
30	customer_value_segment_High value - Low Risk	0.027569
19	usage_level_medium usage	0.018569
16	contract_type_Yearly	0.014822
3	total_charges	0.010580
9	region_North	0.007958
10	region_South	0.007464
5	avg_data_gb	0.006487
11	region_West	0.006439
23	engagement_level_medium engagement	0.005187

[112]:

```
import matplotlib.pyplot as plt

top_features = coef_df.head(10)

plt.figure(figsize=(8,5))
plt.barh(top_features['Feature'], top_features['Coefficient'])
plt.xlabel("Coefficient Value")
plt.title("Top 10 Features Impacting Churn")
plt.gca().invert_yaxis()
plt.show()
```

Top 10 Features Impacting Churn

